

```
# Judge LLM Evaluation Notebook with All Advanced Features
# Requirements: Gemini 1.5 Pro API key, .png chart files, insight text files

import os
import csv
import re
from PIL import Image
import google.generativeai as genai

# === CONFIGURATION ===
# Set your API key and paths for chart images, insight drafts, and output logs
API_KEY = "YOUR_API_KEY" # Replace with your actual Gemini API key
CHART_FOLDER = "charts"
INSIGHT_FOLDER = "insights"
OUTPUT_CSV = "judge_eval_results.csv"
N_VARIANTS = 3 # Number of insight drafts to evaluate per chart for reranking

# === SETUP ===
# Configure Gemini API and instantiate the Gemini 1.5 Pro model
genai.configure(api_key=API_KEY)
model = genai.GenerativeModel("gemini-1.5-pro")

# === PROMPT TEMPLATE ===
# Builds a consistent prompt with rubric instructions for the Judge LLM
# Takes insight text and instructs the model to evaluate across 4 dimensions
# Includes a fail-safe message if the insight violates MTSS alignment rules

def build_prompt(insight_text):
    return f"""
You are an education policy expert reviewing an insight based on a performance chart using a multi-tiered system of supports (MTS)

Evaluate the following:
1. Accuracy
2. Specificity
3. MTSS Alignment (must be actionable, grounded in the chart, no reflection on past)
4. Clarity (≤300 words, ≤30 words/observation)

Insight:
{insight_text}

Please provide:
- Score (1-10) for each dimension
- Brief justification for each
- Suggested rewrites if necessary
If the insight is off-policy or missing, respond: 'No valid insight found.'
"""

# === STRUCTURED OUTPUT PARSING ===
# Extracts scores from the LLM's output response
# Looks for lines containing each rubric name and extracts the numeric score using regex

def parse_response(response_text):
    scores = {"Accuracy": None, "Specificity": None, "MTSS Alignment": None, "Clarity": None}
    for line in response_text.splitlines():
        for key in scores:
            if key in line:
                match = re.search(r'(\d+)', line)
                if match:
                    scores[key] = int(match.group(1))
    return scores

# === EVALUATION FUNCTION ===
# Evaluates a single chart and insight draft using Gemini 1.5 Pro
# Sends the image and text prompt to the model and returns its response

def evaluate_single(chart_path, insight_text):
    image = Image.open(chart_path)
    prompt = build_prompt(insight_text)
    response = model.generate_content([prompt, image])
    return response.text
```

```

# === BATCH RERANKING FUNCTION ===
# Evaluates multiple insight drafts per chart, scores each, and returns the best-scoring one
# Uses structured output parsing to compute average score and handles fail-safe detection

def rank_variants(chart_path, insight_variants):
    best_score = -1
    best_output = None
    best_variant = None
    for i, text in enumerate(insight_variants):
        print(f"\nEvaluating Variant {i+1}")
        raw_response = evaluate_single(chart_path, text)
        print(raw_response)
        if "No valid insight" in raw_response:
            continue # Fail-safe: skip invalid or off-policy responses
        parsed_scores = parse_response(raw_response)
        if None not in parsed_scores.values():
            avg_score = sum(parsed_scores.values()) / 4
            if avg_score > best_score:
                best_score = avg_score
                best_output = raw_response
                best_variant = text
    return best_variant, best_output

# === CSV LOGGING ===
# Appends evaluation results (scores, raw insight, and LLM output) to a CSV file
# Automatically adds headers if the file does not exist yet

def log_to_csv(chart_name, variant_text, scores, full_response):
    header = ["chart", "accuracy", "specificity", "mtss_alignment", "clarity", "insight_text", "judge_output"]
    row = [chart_name, scores.get("Accuracy"), scores.get("Specificity"), scores.get("MTSS Alignment"),
           scores.get("Clarity"), variant_text, full_response]
    file_exists = os.path.exists(OUTPUT_CSV)
    with open(OUTPUT_CSV, 'a') as f:
        writer = csv.writer(f)
        if not file_exists:
            writer.writerow(header)
        writer.writerow(row)

# === MAIN EXECUTION LOOP ===
# Iterates over chart images, evaluates each set of insight variants,
# logs judge feedback to CSV, and saves the best insight to /best_insights/

best_insights_folder = "best_insights"
os.makedirs(best_insights_folder, exist_ok=True)

for filename in os.listdir(CHART_FOLDER):
    if not filename.endswith(".png"):
        continue

    chart_path = os.path.join(CHART_FOLDER, filename)
    base_name = os.path.splitext(filename)[0]
    insight_variants = []

    # Load all insight variants matching this chart
    for i in range(1, N_VARIANTS + 1):
        insight_path = os.path.join(INSIGHT_FOLDER, f"{base_name}_v{i}.txt")
        if os.path.exists(insight_path):
            with open(insight_path, "r") as f:
                insight_variants.append(f.read())

    if not insight_variants:
        print(f"No insight variants found for {base_name}")
        continue

    # Evaluate and rank variants using Judge LLM
    best_text, judge_output = rank_variants(chart_path, insight_variants)
    scores = parse_response(judge_output)

    if scores and best_text:
        # Log to CSV
        log_to_csv(base_name, best_text, scores, judge_output)

        # Export top insight to /best_insights/ as .txt
        output_path = os.path.join(best_insights_folder, f"{base_name}.txt")
        with open(output_path, "w") as f:

```

```

        f.write(best_text)
    print(f"✅ Saved top insight for {base_name} to {output_path}")

# === README.md ===
readme_text = """# 🧑 Judge LLM for Educational Insights (MTSS-Aligned)

This project uses Gemini 1.5 Pro to evaluate AI-generated school insights aligned to multi-tiered system of supports (MTSS). The

- Accuracy
- Specificity
- MTSS Alignment
- Clarity

It also:
- Selects the best variant among multiple insight drafts using average score
- Suggests rewrites if needed
- Filters out off-policy or incomplete insights

---

## Project Structure

```
/judge_llm_project/
└── charts/          # .png chart images
└── insights/         # .txt insight variants per chart
└── best_insights/   # top-ranked insights for dashboard use
└── judge_eval_results.csv # generated output
└── Judge Llm Eval.ipynb # this notebook
```

Each chart must have multiple pre-generated insight variants:
```
chart1_v1.txt
chart1_v2.txt
chart1_v3.txt
```

These must be manually created or generated by a separate script or model *before* running this notebook. This notebook does not

If analyzing multiple schools separately, include the school name in the file prefix to group insights and charts together. Do *

### Example:

**Folder:** `/charts/`
```
schoolA_chart.png
schoolB_chart.png
```

**Folder:** `/insights/`
```
schoolA_chart_v1.txt
schoolA_chart_v2.txt
schoolA_chart_v3.txt

schoolB_chart_v1.txt
schoolB_chart_v2.txt
schoolB_chart_v3.txt
```

**Folder:** `/best_insights/`
```
schoolA_chart.txt # highest-rated insight only
schoolB_chart.txt
```

---

## ✅ Insight Format

Each insight draft should follow this structure:
```
Blurb: A 1-2 sentence summary
```

```

**Observations:**

- Start with percentages or trends
- Avoid interpretation not visible in the chart
- Use ≤ 4 observations, each ≤ 30 words

**Recommendations:**

- School-wide support action
  - Targeted action for at-risk groups
  - Intensive intervention for high-need students
- ```

---

**## 📊 Output**

Logged to `judge\_eval\_results.csv`:

- Chart filename
- Scores for all 4 dimensions
- The original insight
- Full response from Judge LLM

Also saved to `/best\_insights/` for dashboard integration:

- One `.txt` file per chart
- Contains only the top-rated insight body (no judge explanation)

---

**## 🧠 Scoring Rubric**

Dimension	Description
Accuracy	Are the observations grounded in the data?
Specificity	Are numbers or metrics used instead of vague phrases?
MTSS Alignment	Are recommendations immediately actionable, specific, and aligned to MTSS tiers?
Clarity	Is the format readable, within word limits, and educator-friendly?

---

**## 🗂 Evaluation Behavior**

- The judge scores all insight variants for a chart
- It chooses the highest-scoring insight using average score
- If an insight is off-policy, it returns: "No valid insight found."
- Results are logged to CSV
- Top insight text is written to `/best\_insights/` as plain `.txt`

---

**## 🔧 Extending the System**

This notebook currently supports:

- Single-insight evaluation per image
- Reranking up to 3 variants using score average
- Score logging
- Top insight extraction for dashboard display
- Judge prompt for MTSS-aligned insight checking

Planned enhancements:

- Add a second prompt for Judge-generated rewrites of poor outputs
- Comparison prompt: let Judge pick between A vs B
- Prompt for MTSS-specific tier labeling suggestions

---

**## 🚀 How to Use**

1. Generate 2-3 insight variants per chart and save as `.txt` files in `/insights/`
2. Add matching `.png` chart files to `/charts/`
3. Prefix filenames with school names if analyzing multiple schools
4. Replace `YOUR\_API\_KEY` in the notebook
5. Run all cells to score, rank, and log the best insight
6. Read dashboard-ready outputs in `/best\_insights/`

---

**## 🔍 Future Extensions**

- Reward model training (RLAIF)
- Self-revision using Judge feedback
- UI for batch testing or insight generation

```
"""
```

```
# Write README to file
with open("JudgeLLM_README.md", "w") as f:
    f.write(readme_text)
print("JudgeLLM_README.md created.")
```

```
→ JudgeLLM_README.md created.
```